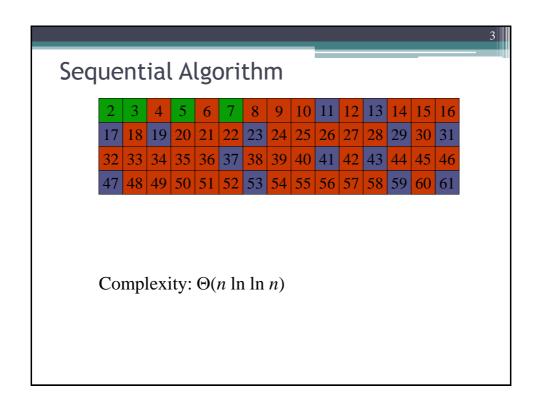
The Sieve of Eratosthenes Jorge Barbosa, FEUP

Outline

- Sequential algorithm
- Sources of parallelism
- Data decomposition options
- $\bullet \ \ {\bf Parallel\ algorithm\ development, analysis}$
- Benchmarking
- Optimizations



Pseudocode

- 1. Create list of unmarked natural numbers 2, 3, ..., n
- $2. k \leftarrow 2$
- 3. Repeat
 - (a) Mark all multiples of k between k^2 and n
 - (b) $k \leftarrow \text{smallest unmarked number} > k$ until $k^2 > n$
- 4. The unmarked numbers are primes

Sources of Parallelism

- Domain decomposition
 - Divide data into pieces
 - Associate computational steps with data
- One primitive task per array element

Making 3(a) Parallel

Mark all multiples of k between k^2 and n



```
for all j where k^2 \le j \le n do
if j \mod k = 0 then
mark j (it is not a prime)
endif
endfor
```

Making 3(b) Parallel

Find smallest unmarked number > k

 \Rightarrow

Min-reduction (to find smallest unmarked number > k)

Shared variable (to get results from all processes)

Agglomeration Goals

- Consolidate tasks
- Reduce sharing costs
- Balance computations among processes

R

Data Decomposition Options

- Interleaved (cyclic)
 - Easy to determine "owner" of each index
 - Leads to load imbalance for this problem
 (with P=2, one processor would be idle after first step)
- Block
 - Balances loads
 - More complicated to determine owner if *n* not a multiple of *p*

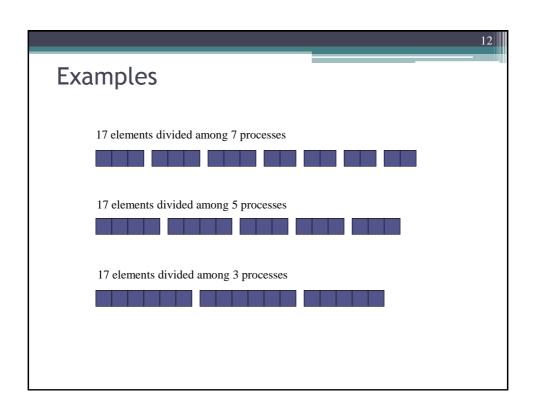
Block Decomposition Options

- Want to balance workload when n not a multiple of p
- Each process gets either $\lceil n/p \rceil$ or $\lfloor n/p \rfloor$ elements
- Seek simple expressions
 - Find low, high indices given an owner
 - Find owner given an index

10

Method #1

- Let $r = n \mod p$
- If r = 0, all blocks have same size
- Else
 - First r blocks have size $\lceil n/p \rceil$
 - Remaining p-r blocks have size $\lfloor n/p \rfloor$



13

Method #1 Calculations

• First element controlled by process *i*

$$i \mid n/p \mid + \min(i,r)$$

• Last element controlled by process i

$$(i+1)\lfloor n/p\rfloor + \min(i+1,r) - 1$$

• Process controlling element *j*

$$\max(\lfloor j/(\lfloor n/p\rfloor+1)\rfloor,\lfloor (j-r)/\lfloor n/p\rfloor)$$

14

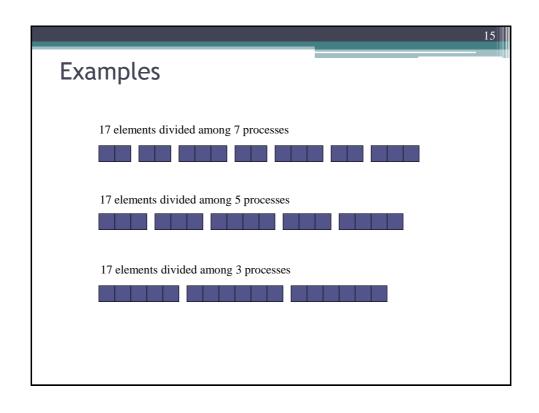
Method #2

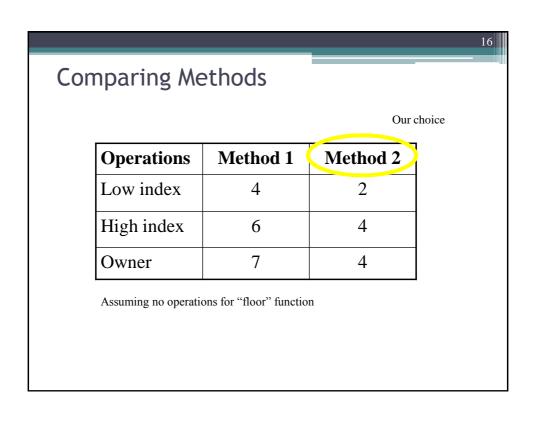
- Scatters larger blocks among processes
- First element controlled by process i |in/p|
- Last element controlled by process i

$$\lfloor (i+1)n/p \rfloor -1$$

- Process controlling element \boldsymbol{j}

$$\lfloor (p(j+1)-1)/n \rfloor$$





17

Block Decomposition Macros

18

Looping over Elements

```
• Sequential program
for (i = 0; i < n; i++) {
    ...
}</pre>
```

Parallel program

```
size = BLOCK_SIZE (id,p,n);
for (i = 0; i < size; i++) {
    gi = i + BLOCK_LOW(id,p,n);
}</pre>
```

19

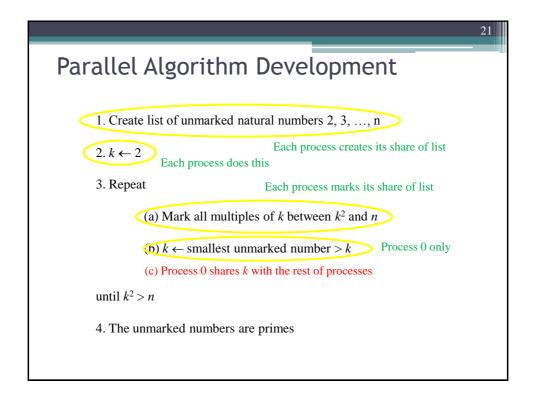
Decomposition Affects Implementation

- Largest prime used to sieve is \sqrt{n}
- First process has $\lfloor n/p \rfloor$ elements
- It has all sieving primes if $p < \sqrt{n}$
- First process always broadcasts next sieving prime
- No reduction step needed

20

Fast Marking

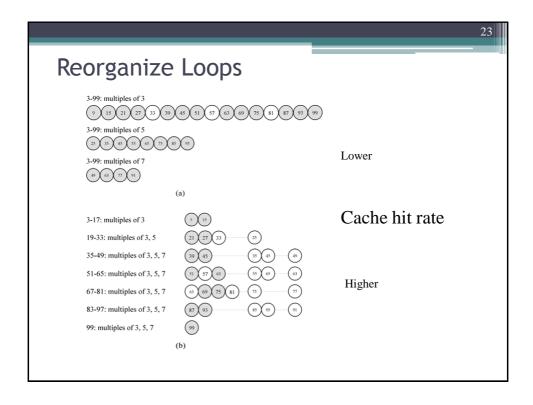
```
Find j the first multiple of k on the block: j, j + k, j + 2k, j + 3k, ... instead of for all j in block if j mod k = 0 then mark j (it is not a prime)
```



Improvements

- Delete even integers
 - Cuts number of computations in half
 - Frees storage for larger values of n
- Each process finds own sieving primes
 - Replicating computation of primes to \sqrt{n}
 - Eliminates broadcast step
- Reorganize loops
 - Increases cache hit rate

22



Lab work

- Develop a shared memory parallelization of the Sieve of Eratosthenes
- Suggestion:
 - Parallel design by domain decomposition
 - Select block distribution
- Consider optimizations to maximize single-processor (core) performance